

Orbit-Evidence

Relational Validity Checks for Learning-Assisted Satellite Communication Experiments

A chronological split constrains order inside one dataset.

Some deployment-validity conditions are not in that dataset.

Anonymous Submission

Ordering alone is not deployment validity

Learning components now sit inside satellite communication software — correcting a propagated orbit, adapting a link parameter, gating a learned branch.

These systems are temporal, so we evaluate with a chronological split. That is *necessary*, and well established.


The claim of this paper

It is **not sufficient** — and the gap is one of *kind*, not degree.

A passing split can coexist with a deployment claim that was never checkable from the data in hand.


Train on the past, test on the future — and still not be able to say the experiment was valid.

Four obligations a split cannot establish



To check any of the four you have to look at something that is not in this table.


The pivot: some validity is relational



Row-local validity is decidable from one realised run. **Relational** validity needs a second execution, a second source state, or a second aggregation level.


Orbit-Evidence: validity as an executable contract

a learning-assisted satellite experiment




Orbit-Evidence: validity as an executable contract


a learning-assisted satellite experiment



Orbit-Evidence: validity as an executable contract



Orbit-Evidence: validity as an executable contract



The two relational checks

L4.6 — differential provenance

Two runs, same *manifest* — the record of everything a run declared it depended on — but **different output** $\Rightarrow$ the record is incomplete.

Refutes a claim of completeness *without enumerating* the dependencies a manifest might omit. One-sided by construction: failure is evidence, success is not certification.


L4.7 — statistical-unit adequacy

Does dependence survive **one declared level coarser**? Null: shuffling the group labels changes nothing.

And it *abstains*. With too few groups, the smallest p the test can *possibly* produce already exceeds α — that design could not have rejected at any effect size. Not weak evidence: no evidence either way. It returns INDETERMINATE.


One judgement generates both: neither obligation can be certified, so we accept checks that can only refuse.

Calibrating the gate before trusting it



On data built with no real dependence the gate still halts 14 times in 450 — against the 0.05 false-alarm budget we allow it. The interval, not the point, is the claim, and it contains that budget.

Does the unit problem appear in real orbital data?



grouping

verdict

pass → element set

$\hat{p} = 0.501$

272 in 90

$p = 0.0025$

lower bound 0.376

HALT

element set → object $\hat{p} = 0.284$

90 in 11


HALT

Both p at $1/400$ — the floor of a 400-permutation test. Measured: how far an orbit solution moves *along its own track* when the next one replaces it — **not truth error**. On a purely *geometric* quantity the same rule returned $\hat{p} = 0.0$: the observable decides.

Can the frozen contract be applied outside this project?

A public spacecraft-telemetry LSTM artifact, frozen at a commit chosen *before* inspection. The detector `sha256` is unchanged throughout — equal to the pre-registration hash.

5 of 19 rules are recorded *not observable* rather than passed. Non-observability is never scored as compliance.



A *disposition* precedes any verdict: the obligation does not arise, or the artifact does not expose the inputs the rule needs. **3 HALT is a verdict on the recorded experiment, not on the published result.**

The contract runs unchanged on an artifact nobody here produced.

Does correcting a detected violation change the experiment?

	original		corrected
overlap	100.0% of val. support	→	0
L4.1	HALT	→	PASS
checkpoint	changed in 5 of 5 paired seeds		

Re-running one seed reproduces its checkpoint **bit for bit**
— the intervention, not training noise.

The *third-party model's* detection metric moved on three seeds, *in both directions*.

Not a null: at 5 paired seeds the smallest attainable two-sided p is $0.0625 > \alpha$.

⇒ **not estimable** at the available paired-run resolution — an evidentiary limit of *this study*, not a rule verdict and not a null result.

Model selection changed; the downstream endpoint is not estimable at this resolution.

What is actually new

Not the primitives.

ICC(1), permutation inference,
metamorphic relations, hermetic build
closure, mutation testing — all established,
none proposed here.

Not a checklist.

The 19 rules are the satellite *instantiation*,
not the contribution.

The method is the conversion

1. Identify which validity obligations are *relational*.
2. Choose the counterfactual execution or coarser level that makes each *falsifiable*.
3. Encode the invariant as a contract — and let an undecidable case return `INDETERMINATE` rather than `pass`.

A permutation reference is ordinarily an analysis step. Here it calibrates a *gate* — with an operating characteristic and an abstention state.

What this does not establish

- **No completeness claim.** L4.6 is bounded by the caller's variation set; the input nobody thought of is absent by construction.
- **No RF or packet-level result.** The object of study is the experiment, not the link.
- **No improvement in learned accuracy.** The contract refuses claims; it does not make models better.
- **Represented-fault regression coverage only.** Faults and rules largely share an origin, so curated coverage measures *reachability*, not fault-finding power.
- **One outside dataset, one outside model.** One step away from marking our own homework — not zero.
- **PASS means not rejected**, never verified.

A validity method must apply the same evidence discipline to its own claims.

**Chronological separation remains necessary.
It is not sufficient.**

Some assumptions must become **executable**, **falsifiable**,
and allowed to **refuse** an unsupported claim.

Backup — why the abstention floor is not a power floor

Take the specific construction of *three coarser groups of two units*: it admits $6!/(2!^3 3!) = 15$ distinct assignments, so the smallest attainable p is $1/15 = 0.067 > \alpha$. That design cannot reject at the nominal level at any effect size.

Four groups of two admit 105 assignments, whose exhaustive reference reaches 0.0095. The general rule is stated at the rule level, not by extrapolating $1/15$: designs below the registered resolution floor return INDETERMINATE.

So INDETERMINATE is a statement about the *design's resolution*, not about the strength of the effect. Collapsing it into PASS is how a design that could not have rejected comes to look like one that did not.

Backup — reproduction and provenance

- make gate regenerates the fault matrix and L4.7's operating characteristic, and *refuses to build* if a headline number disagrees with the artifact at the sentence claiming it. The manuscript contains 64 artifact-bound claim sites; every numerical value shown in this deck is generated from that same summary artifact.
- The two external studies are hash-matched rather than regenerated: their inputs do not ship.
- make external-consequence-verify checks the committed third-party results against the frozen commit, the frozen detector hash and the recorded data hashes — 16 checks, no training.
- The telemetry arrays came from a *checksum-verified mirror* after the upstream endpoint became unavailable; per-file hashes matched two independently published checksum sources. That is mirror concordance, *not* verification by the original publisher.